

# Åland Data Integration Project

## Project Overview

This project involves building a data integration pipeline and data warehouse for analyzing tourism, grocery sales, and demographic data from Åland municipalities. The goal is to create a comprehensive data warehouse that enables analytical insights across multiple data domains.

## Purpose

Build a data pipeline that integrates:

- **Tourism data:** Monthly visitor statistics and revenue by municipality (CSV format)
- **Grocery sales data:** Daily product-level sales transactions from a grocery store chain (JSON format)
- **Demographic data:** Population statistics from ÅSUB (Ålands statistik- och utredningsbyrå) via REST API

The integrated data will be stored in a SQL database using a star schema design, enabling powerful analytical queries and visualizations.

## Data Sources

### 1. Tourism Data (CSV)

- **Location:** `data/tourism/tourism_data.csv`
- **Format:** CSV with monthly granularity
- **Time Period:** January 2000 to December 2025
- **Content:**
  - Municipality-level monthly statistics
  - Visitor counts, accommodation types, origin countries
  - Tourism revenue in EUR
- **Schema:** See `data/DATA_SCHEMA.md` for detailed field descriptions

### 2. Grocery Sales Data (JSON)

- **Location:** `data/grocery/`

- **Format:** Multiple JSON files
  - `stores.json` : Store and municipality reference data
  - `products.json` : Product catalog with categories and pricing
  - `grocery_sales_*.json` : Daily product-level sales transactions (one file per year)
- **Time Period:** January 1, 2000 to December 31, 2025 (daily granularity)
- **Content:**
  - Daily sales at product level
  - Multiple stores across 16 Åland municipalities
  - Sales amounts and units sold
- **Schema:** See [data/DATA\\_SCHEMA.md](#) for detailed field descriptions

### 3. Demographic Data (REST API)

- **Source:** ÅSUB (Ålands statistik- och utredningsbyrå)
- **Endpoint:** [Population Statistics API](#)
- **Content:** Population statistics by municipality, year, age, and gender
- **Time Period:** 1975–2024 (students should fetch relevant years for their analysis)
- **Format:** PX-Web API (students will need to query this REST endpoint)

## Requirements

### Analytical Questions

The data warehouse must be designed to answer the following questions through SQL queries:

- 1. How has sales per capita changed over time?**
  - Requires joining grocery sales, population data, and time dimensions
  - Calculate sales per capita by municipality and time period
- 2. Is there a correlation between sales and tourism statistics?**
  - Analyze relationship between tourism visitor counts/revenue and grocery sales
  - Consider temporal alignment (monthly tourism vs daily sales)
- 3. Which municipalities have the highest sales per capita, and how does this relate to tourism?**
  - Compare sales per capita across municipalities
  - Investigate correlation with tourism metrics
- 4. What are the seasonal patterns in both tourism and grocery sales?**
  - Identify seasonal trends and patterns

- Compare tourism seasonality with sales seasonality

## **5. How do different product categories perform across municipalities?**

- Analyze product category sales by location
- Identify regional preferences or patterns

## **6. Which stores are the top performers, and what factors contribute to their success?**

- Store-level performance analysis
- Investigate factors like location, municipality, size

## **7. How has tourism revenue changed over time by municipality?**

- Time-series analysis of tourism trends
- Municipality-level comparisons

## **8. What is the relationship between population size and total grocery sales?**

- Demographic correlation analysis
- Population growth vs sales growth

## **9. Are there differences in sales patterns between weekdays and weekends?**

- Temporal pattern analysis
- Day-of-week effects on sales

## **10. How do product category sales correlate with tourism seasons?**

- Cross-domain correlation analysis
- Seasonal product preferences during tourism peaks

**Important:** The data warehouse should be designed to answer **all** of the above questions through SQL queries. However, students are only required to create **visualizations for 2 of these questions**. The choice of which 2 questions to visualize is up to the students.

# **Project Goals**

## **Primary Objectives**

### **1. Build a Data Pipeline**

- Extract data from CSV files (tourism data)
- Extract data from JSON files (grocery sales data)
- Extract data from ÅSUB REST API (demographic data)
- Transform and load data into a SQL database

## 2. Implement Data Warehouse in SQL Database

- Choose a SQL database (PostgreSQL, MySQL, SQLite, SQL Server, etc.)
- Implement medallion architecture:
  - **Gold Layer: Required**
    - Final transformed data in star schema format
  - **Bronze Layer: Recommended**
    - Raw data ingestion layer (highly recommended for staging and data lineage)
  - **Silver Layer: Optional**
    - Cleaned and validated data stage (students may choose to implement this for practice, but it is not required)

## 3. Design Star Schema

- Create fact tables for measurable events (sales, tourism)
- Create dimension tables (time, location/municipality, products, stores, demographics)
- Ensure proper relationships and referential integrity

## 4. Create Visualizations

- Build visualizations for 2 selected analytical questions
- Use any visualization tool (Tableau, Power BI, Python/Matplotlib, R, etc.)

# Design Process

**Students should start with a thorough design process before implementation:**

## 1. Data Analysis

- Understand the structure and content of each data source
- Identify relationships and keys for joining data
- Document data quality issues or considerations

## 2. Star Schema Design

- Identify facts (measurable events) and dimensions (descriptive attributes)
- Design fact tables with appropriate grain (level of detail)
- Design dimension tables with all necessary attributes
- Create entity-relationship diagrams
- Consider aggregation strategies

## 3. ETL/ELT Design

- Plan extraction methods for each data source
- Design transformation logic (cleaning, joining, aggregating)

- Plan loading strategy (see Loading Patterns section - full refresh is recommended for this project)
- Consider data refresh frequency

#### 4. Query Design

- Plan SQL queries for each analytical question
- Identify required joins and aggregations
- Consider performance optimization

#### 5. Visualization Planning

- Select 2 questions to visualize
- Design appropriate chart types
- Plan data preparation for visualizations

## Technical Architecture

### Medallion Architecture

The project should follow the medallion architecture pattern:

- **Bronze Layer (Recommended):** Raw, unprocessed data as ingested from sources
  - **Highly recommended** for this project as it provides a staging area for raw data ingestion
  - Allows for data lineage tracking and easier debugging
  - Enables reprocessing of data without re-extracting from source systems
- **Silver Layer (Optional):** Cleaned and validated data, ready for transformation
- **Gold Layer (Required):** Final transformed data in star schema, optimized for analytics

**Note:** The **gold layer** with star schema is **mandatory**. The **bronze layer** is **highly recommended** as it provides a good practice for data ingestion and staging. The silver layer is optional, but students are encouraged to implement it if they want to practice the full medallion architecture.


### Star Schema Design

The gold layer should implement a star schema with:

- **Fact Tables:**
  - Sales fact (grocery sales transactions)
  - Tourism fact (monthly tourism statistics)
  - Potentially combined fact tables depending on design choices
- **Dimension Tables:**

- Time dimension (date, month, quarter, year, day of week, etc.)
- Location/Municipality dimension (municipality codes, names, population data)
- Product dimension (product details, categories, suppliers)
- Store dimension (store information, locations)
- Demographics dimension (population statistics by time and location)

## Data Flow



## Loading Patterns

When loading data into the data warehouse, there are two main approaches:

- **Full Refresh (Truncate and Load)**

- Delete all existing data from target tables
- Load all data from source systems
- Simple to implement and ensures data consistency
- Suitable for smaller datasets or when historical data changes are rare
- May be slower for large datasets

- **Incremental Refresh (Upsert/Merge)**

- Only load new or changed records since the last load
- Requires tracking changes (e.g., timestamps, change data capture)
- More complex to implement but more efficient for large datasets
- Requires careful handling of updates and deletes

**For this project:** A **full refresh** approach is sufficient and recommended. Since this is a learning project with sample data, students should focus on building a working pipeline rather than optimizing for incremental loads. Full refresh is simpler to implement, easier to debug, and adequate for the dataset sizes in this project.

## SQL Database

Students can choose any SQL database that supports:

- Standard SQL syntax

- Foreign key constraints
- Indexing capabilities
- Common data types (dates, numbers, strings)

Recommended options:

- PostgreSQL
- MySQL/MariaDB
- SQLite (for simplicity, though limited for large datasets)
- SQL Server
- Oracle

# Deliverables

## 1. Data Pipeline Code

- Scripts/code for extracting data from all sources
- Transformation logic
- Loading scripts

## 2. Database Schema

- SQL DDL scripts for creating tables
- Documentation of the star schema design
- Entity-relationship diagrams (recommended)

## 3. SQL Queries (OPTIONAL)

- Queries that answer all 10 analytical questions
- Well-documented and organized
- Optional, but good for learning about writing analytical queries

## 4. Visualizations

- 2 visualizations answering selected questions
- Clear titles, labels, and insights

## 5. Documentation (OPTIONAL)

- README explaining the project structure
- Design decisions and rationale
- Instructions for running the pipeline
- Data dictionary
- Optional

# Getting Started

## 1. Review the Data

- Examine the sample data files in `data/` directory
- Read `data/DATA_SCHEMA.md` for detailed schema information
- Explore the ÅSUB API endpoint

## 2. Start with Design

- Don't jump straight into coding
- Create a comprehensive design document
- Design your star schema on paper/diagram first
- Plan your ETL/ELT process

## 3. Implement Incrementally

- Start with one data source
- Build and test your pipeline
- Add additional sources one at a time
- Iterate and refine

## 4. Test Your Queries

- Verify all analytical questions can be answered
- Check data quality and completeness
- Optimize query performance

## 5. Create Visualizations

- Select 2 questions that interest you
- Design clear, informative visualizations
- Ensure visualizations tell a story

# Resources

- **Data Schema Documentation:** `data/DATA_SCHEMA.md`
- **ÅSUB Statistics Portal:** <https://pxweb.asub.ax/>
- **Sample Data:** Located in `data/` directory

# Notes

- The sample data is generated with realistic patterns and correlations (e.g., grocery sales correlate with tourism)
- Municipality data should be fetched from the ÅSUB API, not from the sample data files
- Focus on building a robust, well-designed data warehouse rather than just getting data loaded
- Design decisions should be documented and justified
- Code should be clean, well-commented, and maintainable

Good luck with your project!